

ENGG1811: Computing For Engineers

2025 Term 3

(even more on!) NumPy

Week 8: Monday 3rd November, 2025

Monday 14:00 - 16:00 | HarpM15570

Today

Housekeeping

File-Handling

(even more on) Numpy

Lab Tips

Appendix

Housekeeping

Reminders

- ▶ Assignment 2 is out!
 - ▶ Due 5pm Friday Week 10 (21/11/2025)
 - ▶ Check the course website for details
 - ▶ Same advice as last time:
start early and seek help when needed
- ▶ Self-directed lab 2 on Matlab
 - ▶ Due 5pm Monday Week 11 (24/11/2025)
 - ▶ Very similar to the first self-directed lab: a sequence of videos followed by exercises on MATLAB.
 - ▶ No need to submit ; just complete the exercises on MATLAB Grader; your mark will be the final completion status after the due date.
- ▶ Note these are due around the same time , so don't leave them both to the last minute.

File-Handling

File-Handling: Introduction

- ▶ File-handling is probably the most confusing concept in ENGG1811
 - ▶ Your fundamentals need to be strong
 - ▶ Lists, indexing, for-loops, strings, will all be leveraged when dealing with files

File-Handling: Introduction (Cont)

- ▶ With that being said, files are just really big strings!
 - ▶ Suppose we had a text file with the following riveting content:

```
Hello world 1  
Hello world 2
```

- ▶ Then this is nothing more than:
`"Hello world 1\nHello world 2"`
- ▶ So what are the differences?
 - ▶ A file makes the newline character, `\n`, invisible (we are really inserting this every time we press “enter” on our keyboards!)
 - ▶ Note very carefully, that there is no space between `1`, `\n` and `Hello`.

File-Handling: Files to Strings Examples

- ▶ Lets go through a few examples of this to make sure it's clear:
 - ▶ "Hello Hello Hello 3\n\nHello 1"

Hello Hello Hello 3

Hello 1

- ▶ "Hello"

Hello

- ▶ "\n\n"

File-Handling: Getting Python to Open Files

- ▶ Structure:

```
with open("filename", mode = "r") as f:  
    ## some sort of code here
```

- ▶ `with` is a key-word that does some house-keeping in the background: safely closes the file after you are done using it
- ▶ `open` is a key-word that finds the file named "filename" and converts it into a stream of data for us to use
- ▶ `as f` is a shorthand, just like how we write
`import numpy as np`

- ▶ What is mode doing?

- ▶ There are different modes we can set when accessing a file:
 - ▶ `"r"` - read only
 - ▶ `"w"` - write only
 - ▶ `"a"` - append only
 - ▶ `"r+"` - read and write

File-Handling Function Summary

To summarise, these are the only functions you'll ever need when working with files:

- ▶ `open(file, mode='r')`: opens a file.
- ▶ `np.loadtxt("filename.txt", dtype=str)`: loads the file into a 2D NumPy array (each row is a line, each column is data separated by spaces).
- ▶ `split()`: splits a string (by default on spaces) into a list of words.
- ▶ `readline()`: reads a single line from a file.
- ▶ `readlines()`: reads all lines from a file and returns them as a list of strings.

Examples on the next slides!

File-Handling: Example

- ▶ Suppose we have the following file called `text_file.txt`:

```
This is line 1  
This is line 2
```

- ▶ We access this in Python using:

```
with open("text_file.txt", mode = "r") as f:  
    lines = f.readlines()
```

- ▶ Then `lines` will be the list
`["This is line 1", "This is line 2"]`.
- ▶ The `readlines()` function converts the file into a string, and then will append each line (it knows where a line begins/ends because of the `\n` character).

File-Handling: Example (Cont I)

- ▶ Remember that `lines` is the list `["This is line 1", "This is line 2"]`.
 - ▶ Sometimes getting a list of all the lines might be enough, but what if we want to separate each word from every line?

- ▶ We will access each word using the following for-loop:

```
all_words = []
for i in range(len(lines)):
    words_from_line = lines[i].split()
    all_words.append(words_from_line)
▶ lines[0].split() = ["This", "is", "line", "1"].
▶ lines[1].split() = ["This", "is", "line", "2"].
```

- ▶ **Question:** How do we think `.split()` knows when a word ends?

File-Handling: Example (Cont II)

- We now have

```
all_words = [["This", "is", "line", "1"],  
             ["This", "is", "line", "2"]]
```

We could have also gotten this from

```
all_words = np.loadtxt("text_file.txt", dtype=str)
```

- **Problem:** Python

will assume *everything* is a string/character , even if it's something else!

- it is on us to convert data-types where needed — python will *not* do it for us automatically

- We can do this by the following method:

```
for word_line in all_words:  
    word_line[-1] = int(word_line[-1])
```

- This gives us:

```
all_words = [["This", "is", "line", 1],  
             ["This", "is", "line", 2]]
```

File-Handling: Example (Cont III)

- ▶ Again, suppose we have the following file:

```
This is line 1  
This is line 2
```

- ▶ Using `readline()`, we get one line at a time:

```
line1 = f.readline()  # "This is line 1"  
line2 = f.readline()  # "This is line 2"
```

- ▶ We can then use `split` to get words:

```
# ["This", "is", "line", "1"]  
words_line1 = line1.split()
```

File-Handling: Closing Remarks

- ▶ It is one thing to understand all of this as I go through it, and very different doing this yourself
 - ▶ Difficulty is not in any individual step, but doing/remembering everything in the right order and knowing what to do if (and likely when) you get an error
- ▶ Tips:
 - ▶ If you understand what each operation is doing, less likely for you to forget it or to have a nasty surprise from the output
 - ▶ Do 'all' the file-handling *first*
 - ▶ Process it all , put them nicely into well- labelled lists , and convert data-types
 - ▶ You are much more likely to make an error if you try processing and 'analysing' at the same time!

(even more on) Numpy

NumPy: Shape & Size

1	2	3	4
5	6	7	8
9	10	11	12

All examples are using the top-right array

- ▶ `np.shape(array)`
 - ▶ Gives the shape — the number of rows and columns — of the given array
 - ▶ `np.shape(array) = (3, 4)`
- ▶ `np.size(array)`
 - ▶ Gives the area (number of elements) of the given array
 - ▶ `np.size(array) = 12`

NumPy: Ravel

1	2	3	4
5	6	7	8
9	10	11	12

- ▶ `np.ravel(array)`

Using the array at the top-right

- ▶ Unravels the data into a single horizontal array (list) — converts an n dimensional array into a 1 dimensional array
- ▶ `np.ravel(array) = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]`

- ▶ **Question:** Can we achieve the same result using the reshape function? Why or why not?

NumPy: Slicing

1	2	3	4
5	6	7	8
9	10	11	12

We can slice out a subarray using the format:

```
array[row_start:row_end:row_step,  
      column_start:column_end:column_step]
```

Here are some examples:

- ▶ `array[1:, :2] = array([[5, 6],
 [9, 10]])`
- ▶ `array[::-1, :] = array([[9, 10, 11, 12],
 [5, 6, 7, 8],
 [1, 2, 3, 4]])`
- ▶ `array[:, ::3] = array([[1, 2]])`

NumPy: ix_

1	2	3	4
5	6	7	8
9	10	11	12

- ▶ **Recall:** When we wish to use boolean indexing, we must create an array of the same shape as the array we wish to index, and then fill each entry with either a True or False value
- ▶ For the array at the top right, we might have to make (from scratch) the following array to index it:

```
mask = np.array([[False, False, False, False],  
                 [False, True, True, True],  
                 [False, True, True, True]])
```

- ▶ **Question:** What is the output of array[mask] ?

NumPy: ix_ (Cont I)

1	2	3	4
5	6	7	8
9	10	11	12

- ▶ This is a slightly burdensome procedure — can we be a bit lazier?
 - ▶ Imagine having to create your own boolean array every time you want to do some simple indexing
- ▶ Fortunately, we can!
 - ▶ We have two options: we can select all the rows , and then boolean index on the columns **OR**
we can select all the columns , and boolean index on the rows
 - ▶ **Question I:** What is the output of
`array[[False, True, True], :]` ?
 - ▶ **Question II:** What is the output of
`array[:, [False, True, True]]` ?

NumPy: ix_ (Cont II)

1	2	3	4
5	6	7	8
9	10	11	12

Limitation

What if we want to boolean index both the rows and columns at the same time?

- ▶ What's actually stopping us? Let's try
`array[[False, True, True],
 [False, True, True, True]]`

Question: What should this *intuitively* give us?

NumPy: ix_ (Cont III)

The diagram illustrates the element-wise AND operation between two 2D arrays. On the left, there are two 2x4 arrays. The first array has rows [False, True, True, True] and columns [False, False, False, False]. The second array has rows [False, False, True, True] and columns [True, True, True, True]. A small ampersand (&) symbol is placed between them, indicating the logical AND operation. To the right of the arrays is an equals sign (=), followed by the resulting 2x4 array. This result has rows [False, False, False, False] and columns [False, True, True, True].

False	False	False	False
True	True	True	True
True	True	True	True

&

False	True	True	True
False	True	True	True
False	True	True	True

=

False	False	False	False
False	True	True	True
False	True	True	True

Figure: Intuitive results

- ▶ Annoyingly, we get this back instead:

Error Message From Above

```
IndexError: shape mismatch: indexing arrays could  
not be broadcast together with shapes (2,) (3,)
```

NumPy: ix_ (Cont IV)

1	2	3	4
5	6	7	8
9	10	11	12

- ▶ `np.ix_(rows_boolean_array, columns_boolean_array)`
 - ▶ `rows_boolean_array` is the boolean array which selects which rows you want to keep
 - ▶ `columns_boolean_array` is the boolean array which selects which columns you want to keep
 - ▶ creates a boolean array with desired rows and columns
- ▶ Use case:

```
mask = np.ix_([False, True, True],  
             [False, True, True, True])  
array[mask] = np.array([[6, 7, 8],  
                      [10, 11, 12]])
```

This solves our problem, and we can now be slightly lazier when it comes to boolean indexing (so long as we remember the `ix_` function!)

NumPy: Diff

- ▶ `numpy.diff(array)`
 - ▶ If we have an array $[x_0, x_1, \dots, x_{n-1}]$ then this will return us back the list
$$[x_1 - x_0, x_2 - x_1, \dots, x_{i+1} - x_i, \dots, x_{n-1} - x_{n-2}].$$
 - ▶ In plain English, it is calculating for us the
(forward) difference between consecutive elements of a given
array
- ▶ Example:

```
array = np.array([3, 7, 4, 9, 4, -1])
np.diff(array) = array([ 4, -3, 5, -5, -5])
```

NumPy: Unique

1	1	3	2
2	2	1	3
4	3	4	3

- ▶ `np.unique(array, return_counts = False)`
 - ▶ Returns an array of all the unique values of a given array
 - ▶ `return_counts = False` is an *optional* argument —
 - if it is set to `True`, it will also return a
 - list which gives back the amount of times each unique element occurs
- ▶ Example:

```
unique_vals, counts = np.unique(array,  
                                return_counts = True)  
unique_vals = np.array([1, 2, 3, 4])  
counts      = np.array([3, 3, 4, 2])
```

Lab Tips

Lab Tips

- ▶ Exercise 1:
 - ▶ Convert your numbers from string to float .
 - ▶ Don't average all 5 numbers each time ; the second number in the first line tells you how many to use.
 - ▶ The first 10 files are named `temp0x.txt`, not `tempx.txt`.
- ▶ Exercise 2:
 - ▶ Check your answers on the course website or in the Part A starter code before getting marked off.
 - ▶ No loops allowed in Exercise 2.
 - ▶ Watch for questions that specify a method; you must use that method.
 - ▶ For Question 4, use a different approach to Question 8 from Exercise 2 in Week 7 (i.e. if you used slicing before, use `reshape` here).

Appendix

Appendix Explanation

- ▶ The next slides will explain *broadcasting*, a very elegant NumPy technique.
- ▶ Instead of file handling, there used to be a broadcasting exercise in this week's lab.

NumPy: Broadcasting

- ▶ One of the many neat features of NumPy is that it allows arrays of different sizes to work together
- ▶ Intuitively , what should be the answer of adding these two arrays together?
 - ▶ Example 1:

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|c|} \hline 1 \\ \hline 2 \\ \hline 3 \\ \hline \end{array} = ?$$

- ▶ Example 2:

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline \end{array} = ?$$

NumPy: Broadcasting (Cont I)

Here is what NumPy is *really* doing for the above examples

- ▶ Example 1:

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 3 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline 1 & 1 & 1 & 1 \\ \hline 2 & 2 & 2 & 2 \\ \hline 3 & 3 & 3 & 3 \\ \hline \end{array}$$

- ▶ Example 2:

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 1 & 2 & 3 & 4 \\ \hline 1 & 2 & 3 & 4 \\ \hline \end{array}$$

NumPy: Broadcasting (Cont II)

Let's do a few more examples:

- ▶ Example 3:

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} - \boxed{1} = ?$$

- ▶ Example 4 (perhaps this one won't be intuitive):

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 3 \\ \hline \end{array} \times \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 3 \\ \hline \end{array} = ?$$

NumPy: Broadcasting (Cont III)

► Example 3:

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} - \begin{array}{|c|c|c|c|} \hline 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 1 & 2 & 3 \\ \hline 4 & 5 & 6 & 7 \\ \hline 8 & 9 & 10 & 11 \\ \hline \end{array}$$

► Example 4:

$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 3 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 3 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 2 & 2 & 2 \\ \hline 3 & 3 & 3 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 3 \\ \hline 1 & 2 & 3 \\ \hline 1 & 2 & 3 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 1 & 2 & 3 \\ \hline 2 & 4 & 6 \\ \hline 3 & 6 & 9 \\ \hline \end{array}$$

$(3, 1) \quad (1, 3) \quad (3, 3) \quad (\underline{3}, 3) \quad (3, 3)$

NumPy: Broadcasting (Cont IV)

- ▶ Hopefully from these examples, we can see that broadcasting is only going to work if the arrays are **compatible** :
 1. The two arrays share a dimension of the same size and *,
 2. One of the dimensions — for at least one of the arrays — is one
- ▶ Examples:
 - ▶ 2×3 and 1×3 are compatible — they satisfy both conditions
 - ▶ 5×89 and 5×1 are compatible — they satisfy both conditions
 - ▶ 2×3 and 3×4 are *not* compatible — why?
 - ▶ 5×6 and 5×2 are *not compatible* — why?
- ▶ Question: Why did I put an asterisk over the 'and'? What's the exception?

NumPy: Broadcasting Exercise

If you were curious, this was the exercise that was left out:

- ▶ Two NumPy arrays, `pos` (which has a shape of (6, 2)) and `ref`. One interpretation of the `pos` array is as a list of positions: `pos[0, 0]` and `pos[0, 1]` could store the x and y coordinates of an object, respectively.
- ▶ We ask you to compute the distance between each of the 6 positions in `pos` and the reference position. As a reminder, if we have an arbitrary position (x, y) and reference position (a, b) , then the distance between them is:

$$\sqrt{(x - a)^2 + (y - b)^2}$$

- ▶ Complete this task using functions from the `numpy` library. You are not allowed to use loops. The expected answer and starter code is given on the next slide.

NumPy: Broadcasting Exercise

The expected answer and starter code:

```
import numpy as np

pos = np.array([[ 1.72,    2.56],
                [ 0.24,    5.67],
                [-1.24,    5.45],
                [-3.17,   -0.23],
                [ 1.17,   -1.23],
                [ 1.12,    1.08]]))

ref = np.array([1.22, 1.18])

# The expected answer is approximately:
# [ 1.468,  4.596,  4.928,  4.611,  2.410,  0.141 ]
```

If you attempt it, feel free to get feedback on it!

Feedback

Feel free to provide anonymous feedback about the lab!



Feedback Form