

# ENGG1811: Computing For Engineers

2025 Term 3

(even more on!) NumPy

Week 8: Monday 3<sup>rd</sup> November, 2025

Monday 14:00 - 16:00 | HarpM15570

# Today

Housekeeping

File-Handling

(even more on) Numpy

Lab Tips

Appendix

Housekeeping

# Reminders

- ▶ Assignment 2 is out!
  - ▶ Due 5pm Friday Week 10 (21/11/2025)
  - ▶ Check WebCMS3 for details
  - ▶ Same advice as last time:
    - start early and seek help when needed
- ▶ Self-directed lab 2 on Matlab
  - ▶ Due 5pm Friday Week 10 (21/11/2025)
  - ▶ Very similar to the first self-directed lab: a
    - sequence of videos followed by exercises on MATLAB.
  - ▶ No need to submit ; just complete the exercises on MATLAB Grader; your mark will be the final completion status after the due date.
- ▶ Note these are due on the same day , so don't leave them both to the last minute.

# File-Handling

# File-Handling: Introduction

- ▶ File-handling is probably the most confusing concept in ENGG1811
  - ▶ Your fundamentals need to be strong
  - ▶ Lists, indexing, for-loops, strings, will all be leveraged when dealing with files

## File-Handling: Introduction (Cont)

- ▶ With that being said, files are just really big strings!
  - ▶ Suppose we had a text file with the following riveting content:

```
Hello world 1  
Hello world 2
```

- ▶ Then this is nothing more than:  
"Hello world 1\nHello world 2"
- ▶ So what *are* the differences?
  - ▶ A file makes the newline character, `\n`, invisible (we are really inserting this every time we press “enter” on our keyboards!)
  - ▶ Note very carefully, that there is no space between 1, `\n` and Hello.

## File-Handling: Files to Strings Examples

- ▶ Lets go through a few examples of this to make sure it's clear:

- ▶ `"Hello Hello Hello 3\n\nHello 1"`

```
Hello Hello Hello 3  
  
Hello 1
```

- ▶ `"Hello"`

```
Hello
```

- ▶ `"\n\n"`



# File-Handling: Getting Python to Open Files

## ► Structure:

```
with open("filename", mode = "r") as f:  
    ## some sort of code here
```

- `with` is a key-word that does some house-keeping in the background: safely closes the file after you are done using it
- `open` is a key-word that finds the file named "filename" and converts it into a stream of data for us to use
- `as f` is a shorthand, just like how we write `import numpy as np`

## ► What is `mode` doing?

- There are different modes we can set when accessing a file:
  - `"r"` - read only
  - `"w"` - write only
  - `"a"` - append only
  - `"r+"` - read and write

# File-Handling Function Summary

To summarise, these are the only functions you'll ever need when working with files:

- ▶ `open(file, mode='r')`: opens a file.
- ▶ `np.loadtxt("filename.txt", dtype=str)`: loads the file into a 2D NumPy array (each row is a line, each column is data separated by spaces).
- ▶ `split()`: splits a string (by default on spaces) into a list of words.
- ▶ `readline()`: reads a single line from a file.
- ▶ `readlines()`: reads all lines from a file and returns them as a list of strings.

Examples on the next slides!

## File-Handling: Example

- ▶ Suppose we have the following file called `text_file.txt`:

```
This is line 1  
This is line 2
```

- ▶ We access this in Python using:

```
with open("text_file.txt", mode = "r") as f:  
    lines = f.readlines()
```

- ▶ Then `lines` will be the list `["This is line 1", "This is line 2"]`.
- ▶ The `realines()` function converts the file into a string, and then will append each line (it knows where a line begins/ends because of the `\n` character).

## File-Handling: Example (Cont I)

- ▶ Remember that `lines` is the list `["This is line 1", "This is line 2"]`.
  - ▶ Sometimes getting a list of all the lines might be enough, but what if we want to `separate each word` from every line?
- ▶ We will access each word using the following for-loop:

```
all_words = []  
for i in range(len(lines)):  
    words_from_line = lines[i].split()  
    all_words.append(words_from_line)
```

  - ▶ `lines[0].split() = ["This", "is", "line", "1"]`.
  - ▶ `lines[1].split() = ["This", "is", "line", "2"]`.
- ▶ **Question:** How do we think `.split()` knows when a word ends?

## File-Handling: Example (Cont II)

- ▶ We now have  

```
all_words = [["This", "is", "line", "1"],  
             ["This", "is", "line", "2"]]
```

 We could have also gotten this from  

```
all_words = np.loadtxt("text_file.txt", dtype=str)
```
- ▶ **Problem:** Python  
will assume *everything* is a string/character, even if it's something else!
- ▶ it is on us to convert data-types where needed — python *will not* do it for us automatically
- ▶ We can do this by the following method:  

```
for word_line in all_words:  
    word_line[-1] = int(word_line[-1])
```
- ▶ This gives us:  

```
all_words = [["This", "is", "line", 1],  
             ["This", "is", "line", 2]]
```

## File-Handling: Example (Cont III)

- ▶ Again, suppose we have the following file:

```
This is line 1  
This is line 2
```

- ▶ Using `readline()`, we get one line at a time:

```
line1 = f.readline()  # "This is line 1"  
line2 = f.readline()  # "This is line 2"
```

- ▶ We can then use `split` to get words:

```
# ["This", "is", "line", "1"]  
words_line1 = line1.split()
```

## File-Handling: Closing Remarks

- ▶ It is one thing to understand all of this as I go through it, and very different doing this yourself
  - ▶ Difficulty is not in any individual step, but doing/remembering everything in the right order and knowing what to do if (and likely when) you get an error
- ▶ **Tips:**
  - ▶ If you understand what each operation is doing, less likely for you to forget it or to have a nasty surprise from the output
  - ▶ Do 'all' the file-handling *first*
    - ▶ Process it all , put them nicely into well-labelled lists , and convert data-types
    - ▶ You are much more likely to make an error if you try processing and 'analysing' at the same time!

(even more on) Numpy



# NumPy: Shape & Size

1	2	3	4
5	6	7	8
9	10	11	12

All examples are using the top-right array

- ▶ `np.shape(array)`
  - ▶ Gives the shape — the number of rows and columns — of the given array
  - ▶ `np.shape(array) = (3, 4)`
- ▶ `np.size(array)`
  - ▶ Gives the area (number of elements) of the given array
  - ▶ `np.size(array) = 12`

## NumPy: Ravel

1	2	3	4
5	6	7	8
9	10	11	12

- ▶ `np.ravel(array)`

Using the array at the top-right

- ▶ Unravels the data into a single horizontal array (list) — converts an  $n$  dimensional array into a 1 dimensional array
- ▶ `np.ravel(array) = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]`

- ▶ **Question:** Can we achieve the same result using the reshape function? Why or why not?

## NumPy: Slicing

1	2	3	4
5	6	7	8
9	10	11	12

We can slice out a subarray using the format:

```
array[row_start:row_end:row_step,  
      column_start:column_end:column_step]
```

Here are some examples:

- ▶ `array[1:, :2] = array([[ 5, 6],  
 [ 9, 10]])`
- ▶ `array[::-1, :] = array([[ 9, 10, 11, 12],  
 [ 5, 6, 7, 8],  
 [ 1, 2, 3, 4]])`
- ▶ `array[:, :3, :2:] = array([[1, 2]])`

1	2	3	4
5	6	7	8
9	10	11	12

- **Recall:** When we wish to use boolean indexing, we must create an array of the same shape as the array we wish to index, and then fill each entry with either a `True` or `False` value
- For the array at the top right, we might have to make (from scratch) the following array to index it:  

```
mask = np.array([[False, False, False, False],  
                 [False, True,  True,  True],  
                 [False, True,  True,  True]])
```
- **Question:** What is the output of `array[mask]`?

## NumPy: ix\_ (Cont I)

1	2	3	4
5	6	7	8
9	10	11	12

- ▶ This is a slightly burdensome procedure — can we be a bit lazier?
  - ▶ Imagine having to create your own boolean array every time you want to do some simple indexing
- ▶ Fortunately, we can!
  - ▶ We have two options: we can select all the rows , and then boolean index on the columns **OR** we can select all the columns , and boolean index on the rows
  - ▶ **Question I:** What is the output of `array[[False, True, True], :]` ?
  - ▶ **Question II:** What is the output of `array[:, [False, True, True, True]]`?

## NumPy: ix\_ (Cont II)

1	2	3	4
5	6	7	8
9	10	11	12

### Limitation

What if we want to boolean index both the rows and columns at the same time?

- ▶ What's actually stopping us? Let's try  
`array[[False, True, True],  
 [False, True, True, True]]`

**Question:** What should this *intuitively* give us?

## NumPy: ix\_ (Cont III)

False	False	False	False	&	False	True	True	True	=	False	False	False	False
True	True	True	True		False	True	True	True		False	True	True	True
True	True	True	True		False	True	True	True		False	True	True	True

Figure: Intuitive results

- Annoyingly, we get this back instead:

### Error Message From Above

IndexError: shape mismatch: indexing arrays could not be broadcast together with shapes (2,) (3,)

## NumPy: ix\_ (Cont IV)

1	2	3	4
5	6	7	8
9	10	11	12

- ▶ `np.ix_(rows_boolean_array, columns_boolean_array)`
  - ▶ `rows_boolean_array` is the boolean array which selects which rows you want to keep
  - ▶ `columns_boolean_array` is the boolean array which selects which columns you want to keep
  - ▶ creates a boolean array with desired rows and columns
- ▶ Use case:

```
mask = np.ix_([False, True, True],  
              [False, True, True, True])  
array[mask] = np.array([[6, 7, 8],  
                       [10, 11, 12]])
```

This solves our problem, and we can now be slightly lazier when it comes to boolean indexing (so long as we remember the `ix_` function!)



# NumPy: Diff

- ▶ `numpy.diff(array)`

- ▶ If we have an array  $[x_0, x_1, \dots, x_{n-1}]$  then this will return us back the list

$$[x_1 - x_0, x_2 - x_1, \dots, x_{i+1} - x_i, \dots, x_{n-1} - x_{n-2}].$$

- ▶ In plain English, it is calculating for us the (forward) difference between consecutive elements of a given array

- ▶ Example:

```
array = np.array([3, 7, 4, 9, 4, -1])  
np.diff(array) = array([ 4, -3, 5, -5, -5])
```

## NumPy: Unique

1	1	3	2
2	2	1	3
4	3	4	3

- ▶ `np.unique(array, return_counts = False)`
  - ▶ Returns an array of all the unique values of a given array
  - ▶ `return_counts = False` is an optional argument — if it is set to `True`, it will also return a list which gives back the amount of times each unique element occurs
- ▶ Example:

```
unique_vals, counts = np.unique(array,  
                                return_counts = True)  
unique_vals = np.array([1, 2, 3, 4])  
counts      = np.array([3, 3, 4, 2])
```

## Lab Tips

# Lab Tips

- ▶ Exercise 1:
  - ▶ Convert your numbers from string to float .
  - ▶ Don't average all 5 numbers each time ; the second number in the first line tells you how many to use.
  - ▶ The first 10 files are named `temp0x.txt`, not `tempx.txt`.
- ▶ Exercise 2:
  - ▶ Check your answers on the course website or in the Part A starter code before getting marked off.
  - ▶ No loops allowed in Exercise 2.
  - ▶ Watch for questions that specify a method; you must use that method.
  - ▶ For Question 4, use a different approach to Question 8 from Exercise 2 in Week 7 (i.e. if you used slicing before, use *reshape* here).

## Appendix

## Appendix Explanation

- ▶ The next slides will explain *broadcasting*, a very elegant NumPy technique.
- ▶ Instead of file handling, there used to be a broadcasting exercise in this week's lab.

# NumPy: Broadcasting

- ▶ One of the many neat features of NumPy is that it allows arrays of different sizes to work together
- ▶ Intuitively, what should be the answer of adding these two arrays together?
  - ▶ Example 1:

1	2	3	4
5	6	7	8
9	10	11	12

 + 

1
2
3

 = ?

- ▶ Example 2:

1	2	3	4
5	6	7	8
9	10	11	12

 + 

1	2	3	4
---	---	---	---

 = ?

# NumPy: Broadcasting (Cont I)

Here is what NumPy is *really* doing for the above examples

► Example 1:

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 3 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline 1 & 1 & 1 & 1 \\ \hline 2 & 2 & 2 & 2 \\ \hline 3 & 3 & 3 & 3 \\ \hline \end{array}$$

► Example 2:

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 5 & 6 & 7 & 8 \\ \hline 9 & 10 & 11 & 12 \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline 1 & 2 & 3 & 4 \\ \hline 1 & 2 & 3 & 4 \\ \hline \end{array}$$



## NumPy: Broadcasting (Cont II)

Let's do a few more examples:

► Example 3:

1	2	3	4
5	6	7	8
9	10	11	12

 $-$ 

1
---

 $=$  ?

► Example 4 (perhaps this one won't be intuitive):

1	2	3
---	---	---

 $\times$ 

1
2
3

 $=$  ?

## NumPy: Broadcasting (Cont III)

### ► Example 3:

1	2	3	4
5	6	7	8
9	10	11	12

 $-$ 

1	1	1	1
1	1	1	1
1	1	1	1

 $=$ 

0	1	2	3
4	5	6	7
8	9	10	11

### ► Example 4:

1
2
3

 $\times$ 

1	2	3
---	---	---

 $=$ 

1	1	1
2	2	2
3	3	3

 $\times$ 

1	2	3
1	2	3
1	2	3

 $=$ 

1	2	3
2	4	6
3	6	9

(3,1)      (1,3)      (3,3)      (3,3)      (3,3)

## NumPy: Broadcasting (Cont IV)

- ▶ Hopefully from these examples, we can see that broadcasting is only going to work if the arrays are **compatible** :
  1. The two arrays share a dimension of the same size and \*,
  2. One of the dimensions — for at least one of the arrays — is one
- ▶ **Examples:**
  - ▶  $2 \times 3$  and  $1 \times 3$  are compatible — they satisfy both conditions
  - ▶  $5 \times 89$  and  $5 \times 1$  are compatible — they satisfy both conditions
  - ▶  $2 \times 3$  and  $3 \times 4$  are *not* compatible — why?
  - ▶  $5 \times 6$  and  $5 \times 2$  are *not compatible* — why?
- ▶ **Question:** Why did I put an asterisk over the 'and'? What's the exception?

## NumPy: Broadcasting Exercise

If you were curious, this was the exercise that was left out:

- ▶ Two NumPy arrays, `pos` (which has a shape of  $(6, 2)$ ) and `ref`. One interpretation of the `pos` array is as a list of positions: `pos[0, 0]` and `pos[0, 1]` could store the  $x$  and  $y$  coordinates of an object, respectively.
- ▶ We ask you to compute the distance between each of the 6 positions in `pos` and the reference position. As a reminder, if we have an arbitrary position  $(x, y)$  and reference position  $(a, b)$ , then the distance between them is:

$$\sqrt{(x - a)^2 + (y - b)^2}$$

- ▶ Complete this task using functions from the `numpy` library. You are not allowed to use loops. The expected answer and starter code is given on the next slide.

## NumPy: Broadcasting Exercise

The expected answer and starter code:

```
import numpy as np
```

```
pos = np.array([[ 1.72,  2.56],  
                [ 0.24,  5.67],  
                [-1.24,  5.45],  
                [-3.17, -0.23],  
                [ 1.17, -1.23],  
                [ 1.12,  1.08]])
```

```
ref = np.array([1.22, 1.18])
```

*# The expected answer is approximately:*

*# [ 1.468, 4.596, 4.928, 4.611, 2.410, 0.141 ]*

If you attempt it, feel free to get feedback on it!

## Feedback

Feel free to provide anonymous feedback about the lab!



Feedback Form